


# Convolutional Neural Network

ALS User Meeting

September 11, 2023

Tanny Chavez, Tibbers Hao, Alex Hexemer, Wiebke Koepp, Dylan McReynolds, Eric Roberts, **Zhuowen Zhao**, and Petrus Zwart



# What are convolutional neural networks?

Convolutional neural networks are distinguished from other neural networks by their superior performance with image inputs. They have 3 main types of layers: the convolutional layer, pooling layer, and fully-connected (FC) layer.


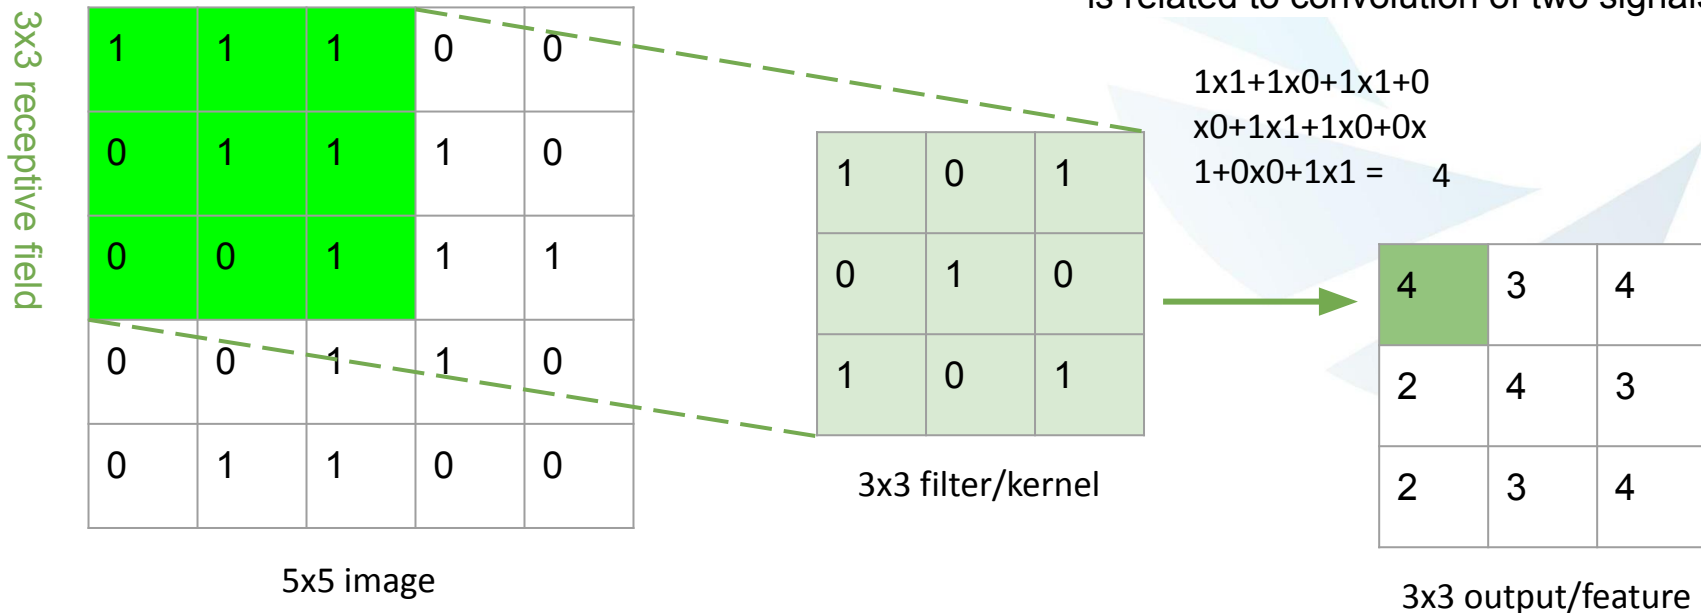


Illustration of LeCun et al. 1998 from CS231n 2017 Lecture 1

# Convolutional layer

The convolutional layer **convolve** the filter with the image i.e. “slide over the image spatially, computing dot products”.

We call the layer convolutional because it is related to convolution of two signals



# Convolutional layer

The convolutional layer **convolve** the filter with the image i.e. “slide over the image spatially, computing dot products”.

|  |   |   |   |   |   |
|--|---|---|---|---|---|
|  | 1 | 1 | 1 | 0 | 0 |
|  | 0 | 1 | 1 | 1 | 0 |
|  | 0 | 0 | 1 | 1 | 1 |
|  | 0 | 0 | 1 | 1 | 0 |
|  | 0 | 1 | 1 | 0 | 0 |

5x5 image

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

3x3 filter/kernel

stride (step size) = 1

output size =  $(N-F)/\text{stride} + 1$

|   |   |   |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| 2 | 3 | 4 |

3x3 output/feature

# Convolutional layer

The convolutional layer **convolve** the filter with the image i.e. “slide over the image spatially, computing dot products”.

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

5x5 image

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

3x3 filter/kernel

|   |   |   |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| 2 | 3 | 4 |

3x3 output/feature

# Convolutional layer

The convolutional layer **convolve** the filter with the image i.e. “slide over the image spatially, computing dot products”.

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

5x5 image

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

3x3 filter/kernel

|   |   |   |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| 2 | 3 | 4 |

3x3 output/feature

# Convolutional layer

The convolutional layer **convolve** the filter with the image i.e. “slide over the image spatially, computing dot products”.

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

5x5 image

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

3x3 filter/kernel

|   |   |   |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| 2 | 3 | 4 |

3x3 output/feature

# Convolutional layer

The convolutional layer **convolve** the filter with the image i.e. “slide over the image spatially, computing dot products”.

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

5x5 image

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

3x3 filter/kernel

|   |   |   |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| 2 | 3 | 4 |

3x3 output/feature



# Convolutional layer

The convolutional layer **convolve** the filter with the image i.e. “slide over the image spatially, computing dot products”.

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

5x5 image

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

3x3 filter/kernel

|   |   |   |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| 2 | 3 | 4 |

3x3 output/feature

# Convolutional layer

The convolutional layer **convolve** the filter with the image i.e. “slide over the image spatially, computing dot products”.

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

5x5 image

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

3x3 filter/kernel

|   |   |   |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| 2 | 3 | 4 |

3x3 output/feature

# Convolutional layer

The convolutional layer **convolve** the filter with the image i.e. “slide over the image spatially, computing dot products”.

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

5x5 image

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

3x3 filter/kernel

|   |   |   |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| 2 | 3 | 4 |

3x3 output/feature

# Padding

When kernel size and stride cannot fit the image dimension, add 0s outside the boundaries.

**padding = 1**

$$\text{output size} = (N-F)/\text{stride} + 1$$

|   |   |   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 1 | 1 | 0 | 1 | 0 | 1 | 0 |
| 0 | 0 | 0 | 1 | 1 | 0 | 0 | 0 | 0 |
| 0 | 1 | 1 | 0 | 0 | 0 | 0 | 1 | 0 |
| 0 | 0 | 0 | 1 | 1 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 0 | 0 | 1 | 1 | 0 | 1 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

7x7 image

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |


3x3 filter/kernel

|   |   |   |
|---|---|---|
| 2 | 2 | 0 |
| 1 | 2 | 1 |
| 0 | 2 | 2 |

3x3 output/feature

# PyTorch implementation

```
torch.nn.Conv2d(in_channels=1, out_channels=1, kernel_size=(3,3), stride=(2,2), padding=1)
```




**Note:** default format for image data in PyTorch: NCHW

# Activation function: ReLU

Adding nonlinearity to the neural network

Pytorch implementation: **torch.nn.ReLU()**



$$\text{ReLU}(x) = (x)^+ = \max(0, x)$$

# Pooling layer

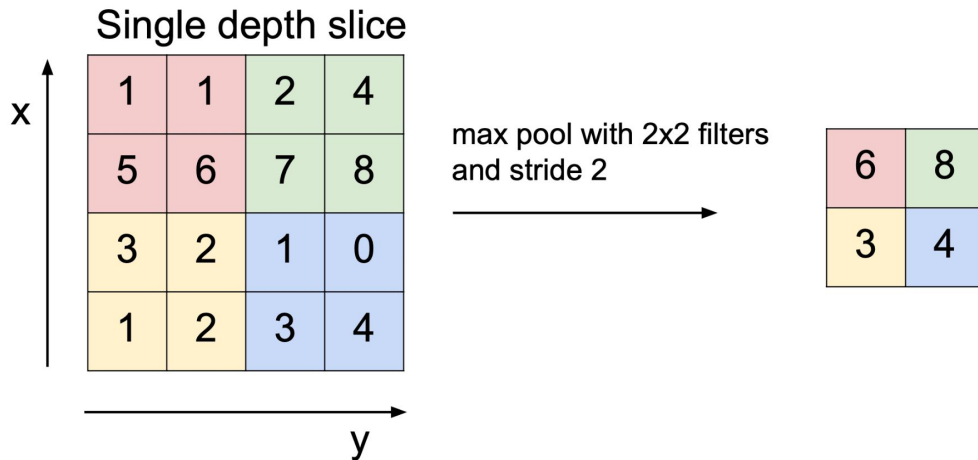
The pooling layer makes the representations smaller and more manageable..

PyTorch implementation: `torch.nn.MaxPool2d(kernel_size=2)`

Other pooling layer APIs: <https://pytorch.org/docs/stable/nn.html#pooling-layers>

## MAX POOLING

- max pooling
- min pooling
- average pooling



# Fully connected layer

In the fully-connected layer, each node in the output layer connects directly to a node in the previous layer.

Pytorch implementation: **`torch.nn.Linear(input_size, output_size)`**

